

Joint Shapley values

Chris Harris¹ Richard Pymar² Colin Rowat³

¹Visual Alpha, Tokyo

²Birkbeck, University of London

³University of Birmingham

23 July 2021

Vector Institute

<https://github.com/harris-chris/joint-shapley-values>



Motivation: joint feature importance

how do I attribute [importance in the presence of] correlations?
(Bhatt et al., 2020)

- FAT (Bhatt et al., 2020)
 - “**local explainability** ...explain the model’s behavior for a specific input”
 - “**global explainability** ...understand the high-level concepts and reasoning used by a model”
- linear models: e.g. $y = \hat{\beta}\mathbf{X} + \varepsilon$ (‘local’ and ‘global’)
- non-linear models?
- **Shapley value**: popular, ‘model agnostic’ approach
 - game theory: average value added by an **individual**, **independently**
 - XAI: average change in model’s prediction due **individual** feature’s value
- problem: feature dependence
 - collinear: individual insignificance (t -test), joint significance (F -test)
 - here: joint feature importance
- contributions
 - uniquely extend Shapley’s value to joint feature importance
 - first index to do so

- 1 Cooperative games
- 2 Joint Shapley values
- 3 Interaction indices
- 4 AI/ML attribution problem
- 5 References

Games in characteristic function form

- most studied class of cooperative game (von Neumann and Morgenstern, 1944)
- set of **agents**, $i \in N = \{1, 2, \dots, n\}$
- value assigned to coalitions, $S \subseteq N$ by **value function**, $v : 2^N \rightarrow \mathbb{R}$
- **Shapley (1953) value**: what's a 'fair' split of $v(N)$?

$$\phi_i^S(v) \equiv \sum_{S \subseteq N \setminus \{i\}} p^i(S) [v(S \cup \{i\}) - v(S)]$$

where $p^i(S) = \frac{|S|!(n-|S|-1)!}{n!}$ randomises uniformly over singletons

- sum of value i that adds to possible coalitions S
- weighted by the (symmetric, independent) probability that i joins S
- **arrival order**: S arrives (any order), then i , then $N \setminus (S \cup i)$ (any order)
- ϕ^S **uniquely** satisfies ① **efficiency**, ② **null/dummy**, ③ **symmetry/anonymity**, ④ **additivity/linearity** axioms

Example (The $n = 3$ glove/market game e.g. Lucas (1971))

Let $i = 1$ have a left glove, and $i \in \{2, 3\}$ each have a right glove, with unit value arising from a pair, so that:

$$v(\emptyset) = v(1) = v(2) = v(3) = v(2, 3) = 0$$

$$v(N) = v(1, 2) = v(1, 3) = 1$$

The Shapley values are:

$$\phi_1^S(v) = \frac{2}{3} > \phi_2^S(v) = \phi_3^S(v) = \frac{1}{6}.$$

Consistent with intuition, the Shapley value privileges agent 1. Less intuitively, it gives no sign that value arises from particular pairs.

Extending Shapley from singletons to sets

- what's the average value a set adds?
- let agents arrive as sets, T (including singletons):

$$\phi_T^J(v) \equiv \sum_{S \subseteq N \setminus T} p^T(S) [v(S \cup T) - v(S)]$$

- extend $p^i(S)$ to $p^T(S)$ by randomising uniformly over sets
- S arrives (any weak order), then T together, then $N \setminus (S \cup T)$ (any weak order)
- will add an **order of explanation**, k , to efficiency axiom
 - controls computational costs, $\mathcal{O}(3^n \wedge (2^n n^k))$
 - introduced by Dhamdhere, Agarwal, and Sundararajan (2020)
 - $k = 1$ reduces to original Shapley
- two branches of game theory literature:
 - 1 fix coalitions *a priori* (Owen, 1977)
 - 2 decompose sets recursively to singletons (Grabisch and Roubens, 1999)

Extending Shapley's axioms: the easy part

(JEF) **joint efficiency**: total worth partitions among sets

$$\sum_{\substack{T \subseteq N \\ |T| \leq k}} \phi_T^J(v) = v(N)$$

(JNU) **joint null**: sets adding no worth get no value

$$v(S \cup T) = v(S) \forall S \subseteq N \setminus T \Rightarrow \phi_T^J(v) = 0$$

(n.b. **interaction indices** build on singletons)

(JLI) **joint linearity**:

$$\phi_T^J(\alpha u + \beta v) = \alpha \phi_T^J(u) + \beta \phi_T^J(v)$$

for any non-negative scalars α and β .

(helps extend proofs from particular games to all games)

Extending Shapley's axioms: the harder part

Shapley's symmetry: any arrival order of agents is equally likely

(ANO) **anonymity** for all permutations, σ , on N ,

$$\psi_i(v) = \psi_{\sigma(i)}(\sigma v)$$

for all $i \in N$, all games v and all $\psi : N \mapsto \mathbb{R}$.

(SYM) **symmetry**

$$v(S \cup i) = v(S \cup j) \forall S \subseteq N \setminus \{i, j\} \Rightarrow \phi_i^S(v) = \phi_j^S(v)$$

- **ANO** \Rightarrow **SYM** (Malawski, 2020)
- Shapley's value uniquely solves **EFF**, **DUM**, **LIN** and either **ANO**/**SYM**
- what's the right way to generalise these? (There are lots of wrong ways)

Joint anonymity and symmetry

(JAN) **joint anonymity**: for all permutations, σ , on N ,

$$\psi_T(v) = \psi_{\sigma(T)}(\sigma v)$$

for all $T \subseteq N$, all games v and all $\psi : 2^N \mapsto \mathbb{R}$

(JSY) **joint symmetry**: if two sets add equal worth to sets that they can both join **and** add no worth to all other coalitions, they receive an equal value:

$$v(S \cup T) = v(S \cup T') \text{ for all } S \subseteq N \setminus (T \cup T')$$

$$v(S \cup T) = v(S) \text{ for all } S \subseteq N \setminus T \text{ s.t. } S \cap T' \neq \emptyset$$

$$v(S \cup T') = v(S) \text{ for all } S \subseteq N \setminus T' \text{ s.t. } S \cap T \neq \emptyset$$

$$\Rightarrow \phi_T(v) = \phi_{T'}(v)$$

Joint Shapley values are unique for each k

Theorem

For each order of explanation $k \in \{1, \dots, n\}$, there is a **unique** ϕ^J which satisfies axioms **JLI**, **JNU**, **JEF**, **JAN** and **JSY**:

$$\phi_T^J(v) = \sum_{S \subseteq N \setminus T} q_{|S|} [v(S \cup T) - v(S)]$$

for each $\emptyset \neq T \subseteq N$ with $|T| \leq k$, where (q_0, \dots, q_{n-1}) uniquely solves

$$q_0 = \frac{1}{\sum_{i=1}^k \binom{n}{i}}, \quad q_r = \frac{\sum_{s=(r-k) \vee 0}^{r-1} \binom{r}{s} q_s}{\sum_{s=1}^{k \wedge (n-r)} \binom{n-r}{s}};$$

for all $r \in 1, \dots, n-1$.

Arrival orders, probabilities are independent of coalition size: $p^T(S) = q_{|S|}$

Proof sketch: if ϕ satisfies ...

- ① **JLI, JNU** \Rightarrow there exist constants $\{p^T(S)\}$ such that ...

$$\phi_T(v) = \sum_{S \subseteq N \setminus T} p^T(S) [v(S \cup T) - v(S)].$$

\therefore no discrete derivatives, unlike Grabisch and Roubens (1999) etc.

- ② **JLI, JNU, JEF** \Leftrightarrow for each order of explanation k ,

$$\delta_N(S) = \sum_{\substack{\emptyset \neq T \subseteq S: \\ |T| \leq k}} p^T(S \setminus T) - \sum_{\substack{\emptyset \neq T \subseteq N \setminus S: \\ |T| \leq k}} p^T(S),$$

for all $\emptyset \neq S \subseteq N$, where $\delta_N(S)$ equals 1 if $S = N$ and 0 otherwise.

- ③ **JLI, JNU, JEF, JAN** \Leftrightarrow

$$p^T(S) = p^{T'}(S') \quad \forall S \subseteq N \setminus T, S' \subseteq N \setminus T' \text{ s.t. } |S| = |S'|, |T| = |T'|$$

- ④ **JLI, JNU, JEF, JSY** $\Leftrightarrow p^T(S) = p^{T'}(S) \quad \forall S \subseteq N \setminus (T \cup T')$

\therefore **JAN, JSY** not nested; non-existence w/o extra **JSY** terms

Interaction indices

- interaction indices assess interactions **within** sets
- joint Shapley measures the value-added **of** a set of features
- ① **Shapley interaction index**, ϕ^{SI} (Grabisch and Roubens, 1999)
 *i and j ...exhibit a positive interaction when the worth of coalition $\{i, j\}$ is more than the sum of **individual worths** ...*

$$v(S \cup \{i, j\}) - v(S \cup \{i\}) - v(S \cup \{j\}) + v(S)$$

- ② **added-value index**, ϕ^{AV} (Alshebli et al., 2019)
*the difference between the outcome of that group and the expected contribution of **each member***
- ③ **Shapley-Taylor**, ϕ^{ST} (Dhamdhere, Agarwal, and Sundararajan, 2020)
interactions of subsets up to some size k ...analogous to how the truncated Taylor series decomposes the function value

Example ($n = 3$ glove game)

T	ϕ_i	ϕ^{SI}	ϕ^{AV}	$\phi^{ST}(2)$	$\phi^{ST}(3)$	$\phi^J(2)$	$\phi^J(3)$
1	2/3	2/3	-5/12	0	0	7/18	7/21
2	1/6	1/6	-1/6	0	0	1/18	1/21
1, 2		1/2	5/12	2/3	1	4/18	4/21
2, 3		-1/2	-1/3	-1/3	0	1/18	1/21
N		-1	1/4		-1		3/21

Example ($n = 3$ majority game: $v(i) = 0 < v(j, k) = v(N) = 1$)

T	ϕ_i	ϕ^{SI}	ϕ^{AV}	$\phi^{ST}(2)$	$\phi^{ST}(3)$	$\phi^J(2)$	$\phi^J(3)$
i	1/3	1/3	-1/3	0	0	1/9	2/21
i, j		0	1/3	1/3	1	2/9	4/21
N		-2	0		-2		3/21

Example ($n = 2$ collinearity game: $v(1) = v(2) = v(N) = 1$)

T	ϕ_i	ϕ^{SI}	ϕ^{AV}	$\phi^{ST}(2)$	$\phi^J(2)$
i	1/2	1/2	1/4	1	1/3
N		-1	-1/2	-1	1/3

Shapley & explainable AI

- in large, highly non-linear models, how explain a feature's importance?
- Lipovetsky and Conklin (2001), Štrumbelj and Kononenko (2014): use Shapley's value
 - agents \Rightarrow features
 - characteristic function $v \Rightarrow$ prediction function f
 - $i \in S \Rightarrow$ evaluated feature i at specific x_i
 - $i \in N \setminus S \Rightarrow$ evaluate feature at mean \bar{x}_i
 - ϕ_i : how much does a specific value of x_i change the prediction?
- implementational decisions include: condition means on observational (q.v. Lundberg and S.-I. Lee, 2017; Frye, Mijolla, et al., 2021) or interventional (q.v. Datta, Sen, and Zick, 2016; Janzing, Minorics, and Blöbaum, 2020; Sundararajan and Najmi, 2020) data?
- for simplicity, and to ease replicability, we call the popular SHAP package (Lundberg and S.-I. Lee, 2017; Lundberg, Erion, et al., 2020)

Boston housing data (Harrison and Rubinfeld, 1978)

- 12 numerical features, and one binary feature; 506 data points
- Dhamdhere, Agarwal, and Sundararajan: random forest regression



Shapley ($k = 1$)	$k = 2$	$\phi^{ST}(f; k)$ $k = 3$	$k = 2$	$\phi^J(f; k)$ $k = 3$
RM: 2.57	RM: 3.12	RM: 3.12	LSTAT: 0.63	LSTAT: 0.42
LSTAT: 2.47	LSTAT: 2.04	LSTAT: 2.04	RM: 0.56	RM: 0.34
AGE: 0.81	DIS: 1.55	DIS: 1.55	LSTAT, RM: 0.30	AGE: 0.11
DIS: 0.63	CRIM: 1.37	CRIM: 1.37	AGE, LSTAT: 0.21	LSTAT, RM: 0.11
CRIM: 0.46	B: 1.31	DIS, LSTAT: 1.33	AGE, RM: 0.20	NOX: 0.10
NOX: 0.44	NOX: 1.15	B: 1.31	DIS, RM: 0.19	DIS: 0.08
PTRATIO: 0.27				
B: 0.24	CHAS, RM: 0.15	AGE, CRIM, RM: 0.21	CHAS, TAX: 0.02	RAD, TAX, ZN: 0.00
TAX: 0.22	LSTAT, TAX: 0.15	AGE, DIS, PTRATIO: 0.21	RAD, ZN: 0.01	CHAS: 0.00
INDUS: 0.19	INDUS, RAD: 0.15	DIS, LSTAT, PTRATIO: 0.21	CHAS, RAD: 0.01	CHAS, RAD, TAX: 0.00
RAD: 0.13	NOX, TAX: 0.14	AGE, LSTAT, PTRATIO: 0.20	ZN: 0.01	CHAS, RAD, ZN: 0.00
ZN: 0.07	DIS, INDUS: 0.13	AGE, NOX, RM: 0.20	CHAS: 0.01	CHAS, TAX, ZN: 0.00
CHAS: 0.07	DIS, PTRATIO: 0.12	B, CRIM, LSTAT: 0.19	CHAS, ZN: 0.01	CHAS, ZN: 0.00

- $k = 2$: ϕ^J singletons, pairs mix evenly: JNU doesn't favour singletons
- $k = 3$: ϕ^J triples small, but shift ranking of singletons, pairs

Movie reviews (Pang and L. Lee, 2005)

- 10,600 binary reviews (100 test); encoded as 1,004-vector (BoW)
- fully connected NN (2 hidden layers, 16 units/layer, ReLU activations)


Review	joint Shapleys
1. negation : aficionados of the whodunit won't be disappointed	$\{\text{disappointed}\}: -2 \times 10^{-5}$ $\{\text{won't}\}: 6 \times 10^{-5}$ $\{\text{be, disappointed}\}: -9 \times 10^{-8}$ $\{\text{won't, disappointed}\}: +6 \times 10^{-8}$ $\{\text{won't, be, disappointed}\}: 5 \times 10^{-9}$
2. enhancement : both inspiring and pure joy	$\{\text{both}\}: 2 \times 10^{-4}$ $\{\text{and}\}: 6 \times 10^{-5}$ $\{\text{and, both}\}: 1 \times 10^{-6}$
3. context : you wish Jacquot had left well enough alone	$\{\text{you, well}\}: +9 \times 10^{-7}$ $\{\text{left, well}\}: -3 \times 10^{-7}$ $\{\text{would}\}: -1 \times 10^{-4}$
4. lost potential : fascinating little thriller that would have been perfect	$\{\text{fascinating}\}: 2 \times 10^{-4}$ $\{\text{would, fascinating}\}: +3 \times 10^{-7}$ $\{\text{would, been, fascinating}\}: -1 \times 10^{-8}$ $\{\text{effort}\}: -1 \times 10^{-5}$
5. qualifying adjective : director ...award-winning ...make a terrific effort	$\{\text{director}\}: -9 \times 10^{-6}$ $\{\text{terrific, effort}\}: +8 \times 10^{-7}$ $\{\text{winning, director}\}: +5 \times 10^{-7}$

- joint Shapleys have a direct interpretation
- BoW understates pairs, triples: co-occurrence rather than k -gram

Sampling joint Shapley values



Sampled ϕ^J converges to exact ϕ^J ;
 $k = 2$ Boston



Difference between consecutive ϕ^J
 samples averages converges to zero;
 $k = 2$ movie review #2

Conclusion, discussion

- direct extension of Shapley value, from singletons to sets
 - how much value does a set of agents add?
 - how much does a set of feature change predictions?
 - intuitive, direct interpretation
 - existing interaction indices assess value-added within sets
- arrival order interpretation allows incorporation of causal knowledge (Frye, Rowat, and Feige, 2020)

Example ($n = 3$ majority game, glove game)

Let $i = 1$ arrive first (causal ancestor).

T	ϕ_i	ϕ^{SI}	ϕ^{AV}	$\phi^{ST}(2)$	$\phi^J(1)$	$\phi^J(2)$
1	?	?	?	?	0	0
2; 3		?	?	?	1/2	1/3
2, 3		?	?	?		1/3

$\phi^J(2)$ indicates that value only accrues to **any** agent arriving second.

References I



Alshebli, B. K., T. P. Michalak, O. Skibski, M. Wooldridge, and T. Rahwan (2019). “A measure of added value in groups”. *ACM Transactions on Autonomous and Adaptive Systems (TAAS)* 13.4, pp. 1–46.



Bhatt, U., A. Xiang, S. Sharma, A. Weller, A. Taly, Y. Jia, J. Ghosh, R. Puri, J. M. F. Moura, and P. Eckersley (2020). “Explainable machine learning in deployment”. In: *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pp. 648–657.



Datta, A., S. Sen, and Y. Zick (2016). “Algorithmic Transparency via Quantitative Input Influence: Theory and Experiments with Learning Systems”. In: *2016 IEEE Symposium on Security and Privacy (SP)*, pp. 598–617.



Dhamdhere, K., A. Agarwal, and M. Sundararajan (2020). “The Shapley Taylor Interaction Index”. In: *International Conference on Machine Learning*. PMLR, pp. 9259–9268.



Frye, C., D. de Mijolla, T. Begley, L. Cowton, M. Stanley, and I. Feige (2021). “Shapley explainability on the data manifold”. In: *International Conference on Learning Representations*.



Frye, C., C. Rowat, and I. Feige (2020). “Asymmetric Shapley values: incorporating causal knowledge into model-agnostic explainability”. *Advances in Neural Information Processing Systems* 33.



Grabisch, M. and M. Roubens (1999). “An axiomatic approach to the concept of interaction among players in cooperative games”. *International Journal of Game Theory* 28.4, pp. 547–565.



Harrison, D. and D. L. Rubinfeld (1978). “Hedonic housing prices and the demand for clean air”. *Journal of environmental economics and management* 5.1, pp. 81–102.



Janzing, D., L. Minorics, and P. Blöbaum (26–28 Aug 2020). “Feature relevance quantification in explainable AI: A causal problem”. In: *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*. Ed. by S. Chiappa and R. Calandra. Vol. 108. Proceedings of Machine Learning Research. PMLR, pp. 2907–2916.

References II



Lipovetsky, S. and M. Conklin (2001). "Analysis of regression in game theory approach". *Applied Stochastic Models in Business and Industry* 17.4, pp. 319–330.



Lucas, W. F. (Oct. 1971). "Some Recent Developments in n -person game theory". *SIAM Review* 13.4, pp. 491–523.



Lundberg, S. M., G. Erion, H. Chen, A. DeGrave, J. M. Prutkin, B. Nair, R. Katz, J. Himmelfarb, N. Bansal, and S.-I. Lee (2020). "From local explanations to global understanding with explainable AI for trees". *Nature machine intelligence* 2.1, pp. 2522–5839.



Lundberg, S. M. and S.-I. Lee (2017). "A unified approach to interpreting model predictions". In: *Advances in Neural Information Processing Systems*, pp. 4765–4774.



Malawski, M. (2020). "A note on equal treatment and symmetry of values". In: *Transactions on Computational Collective Intelligence XXXV*. Springer, pp. 76–84.



Owen, G. (1977). "Values of games with *a priori* unions". In: *Mathematical economics and game theory. Essays in honor of Oskar Morgenstern*. Ed. by R. Henn and O. Moeschlin. Lecture Notes in Economics and Mathematical Sciences 141. Springer, pp. 76–88.



Pang, B. and L. Lee (2005). "Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales". In: *Proceedings of the ACL*.



Shapley, L. S. (1953). "A value for n -person games". In: *Contributions to the theory of games*. Ed. by H. W. Kuhn and A. W. Tucker. Vol. II. Annals of Mathematical Studies 28. Princeton: Princeton University Press. Chap. 17, pp. 307–317.



Štrumbelj, E. and I. Kononenko (2014). "Explaining prediction models and individual predictions with feature contributions". *Knowledge and Information Systems* 41.3, pp. 647–665.

References III



Sundararajan, M. and A. Najmi (13–18 Jul 2020). “The Many Shapley Values for Model Explanation”. In: *Proceedings of the 37th International Conference on Machine Learning*. Ed. by H. D. III and A. Singh. Vol. 119. Proceedings of Machine Learning Research. PMLR, pp. 9269–9278.



von Neumann, J. and O. Morgenstern (1944). *Theory of Games and Economic Behavior*. 1st edition. Princeton University Press.